


# Project 4: Machine Learning

*Team #7: Michael Herold, Ralph Samer*



# Recap

## Task

- Automatic tagging of unseen posts using both supervised & unsupervised approach

## Data set

- StackExchange Data Dump from March, '16
- Subgroup: [programmers.stackexchange.com](http://programmers.stackexchange.com)
- Training set: 34,500 posts (~90%)
- Test set: 3,800 posts (~10%)

## StackExchange

### Total number of ...

|       |        |
|-------|--------|
| Posts | 38,315 |
| Tags  | 1,614  |

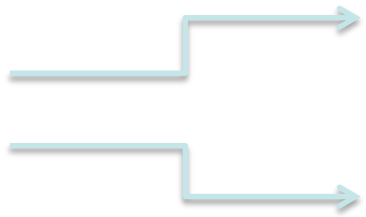
### Tags per post

|         |      |
|---------|------|
| Min     | 1    |
| Max     | 5    |
| Average | 2.68 |

## Current Approach

# Preprocessing

### Posts

- Tokenized content
  - Ignored posts
- 
- *Body*
  - *Title*
  - *Accepted answer (if not negative score)*
  - *Score < threshold and no accepted answer*
  - *No answers and negative score*
- Removed HTML, code, URLs, emoticons, numbers, single letters
  - Removed stop words
  - Stemming (*nltk.stem.porter.PorterStemmer*)
  - Merge synonym words

### Tags

- Merge similar tags
- Remove less frequent tags

## Current Approach

# Features & Transformation

**Feature Selection:**

- word occurrences as features
  - remove words having little impact on performance of our models
- Try additional features that improve performance  
(e.g. whether tag appears in title or body of a document/post or not)

**Challenges:**

- highly depends on preprocessing (stop-word removal, merge synonyms, ...)
- crucial for performance (of models)
- when words are features:
  - important words should appear frequently
  - but not in every single document/post

**Goal:**  
get best prediction accuracy

**Transformation:**

- e.g. TF-IDF, less frequent words occurring more often  
in one post are more important for that post than others  $TF \times IDF_{ij} = TF_{ij} \log\left(\frac{N}{DF_j}\right)$

## Current Approach

# Supervised (I)

Start to train a Multinomial Naive Bayes classifier (*MultinomialNB* from *sklearn*)

### Two Naive Bayes approaches:

- First approach:
  - Train multiple *binary* classifiers (1 classifier/tag)
  - classes = {true, false}
  - Apply classifier on test set => suggest k most probable tags
- Second approach:
  - Train a single classifier for all  $n$  tags
  - #classes =  $n$
  - Apply classifier on test set => suggest k most probable tags

### Reasons for choosing classification over (linear) regression:

- We have  $n$  Tags:  $n$  is a discrete/natural number
- Also: predicting if a post contains a specific tag is discrete/binary problem
- Predicting tags is a typical classification problem
  - tags act as classes (= output of classifiers)

## Current Approach

# Supervised (II)

### **First tests with Naïve Bayes** (full data set):

- sufficient results, but still enough space left for improvements (e.g. improve preprocessing, feature selection, weight important features, ...)

Compare performance of Naïve Bayes with other supervised models:

- **k-Nearest Neighbors:**

- suggest  $h$  most probable tags
- probability depends on:
  - number of neighbors that contain these suggested tags
  - distance to these neighboring documents/posts


- **Linear SVM:**

- Use only a few features that represent the entire post

## Current Approach

# Unsupervised

- **Approach:**
  - Cluster similar posts according to the *TFxIDF* matrix
  - Assign  $n$  most frequent tags of cluster to new samples of cluster
- **k-Means clustering:**
  - Number of clusters (k): #tags
  - Initialization: *k-means++*
  - Distance: *Euclidean*
- **Hierarchical Agglomerative Clustering:**
  - Linkage: *Ward criterion* (minimum variance)
  - Distance: *Euclidean*
  - Stop criterion: number of clusters (#tags)



# Evaluation

- $Precision = \frac{true\_positives}{true\_positives + false\_positives}$

- $Recall = \frac{true\_positives}{true\_positives + false\_negatives}$



Both should be high → **F1 measure**

- True positives: if tag found in prediction set and original set
- False positives: if tag found in prediction set but not in original set
- False negatives: if tag found in original set but not in prediction set



## Current Approach

# First results

| Model                 | Precision | Recall |
|-----------------------|-----------|--------|
| MultinomialNB         | 0.332     | 0.285  |
| k-Means <sup>1)</sup> | 0.298     | 0.262  |
| HAC <sup>2)</sup>     | 0.281     | 0.255  |

All results are based on two predicted tags (average tag assignment = 2.68)

<sup>1)</sup> Number of clusters = #tags

<sup>2)</sup> Stop criterion = #clusters → #tags

# Problems

## Preprocessing:

- High dimensionality → sparse matrix
- POS-tagging and lemmatization not working good enough
- Merge synonyms and similar tags

## Learning:

- Choosing proper number of clusters ( $k$ ) → dendrogram from HAC
- Finding best fitting model and #tags to predict

## Runtime & memory:

- Too many features for forward/backward selection
- HAC very intensive to compute

# Expected results & Outlook

## Expected results:

- F1 not ideal but tags are reasonable
- Precision decreases by the number of suggested tags
- Recall increases by the number of suggested tags

## Outlook:

- Improve preprocessing
- Use different (more independent) features for Naïves Bayes
- Evaluation

# Thank you for your attention.

## References:

- *Knowledge Discovery and Data Mining 1* by Denis Helic and Roman Kern (2015)
- Segaran, T. 2007. *Programming Collective Intelligence: Building Smart Web 2.0 Applications*. Sebastopol: O'Really